


Workshop 7 - Interpretability with Class Activation Mapping



Antonio Rueda-Toicen

About me

- AI Researcher at [Hasso Plattner Institute](#), AI Engineer & DevRel for [Voxel51](#)
- Organizer of the [Berlin Computer Vision Group](#)
- Instructor at [Nvidia's Deep Learning Institute](#) and Berlin's [Data Science Retreat](#)
- Preparing a [MOOC for OpenHPI](#) (now open for registration)



[LinkedIn](#)

How to use our Discord channel during the workshop

- Our channel is **#practical-computer-vision-workshops**. Please ask all questions there instead of the Zoom chat. Through Discord we can have better and more detailed discussions.
- **Step 1 - Use the Discord invite on the [Voxel51 website](https://discord.com/invite/fiftyone-community)**
<https://discord.com/invite/fiftyone-community>
- **Step 2 - Access channel** (use the direct link below or search)
<http://bit.ly/3YmvPXG>

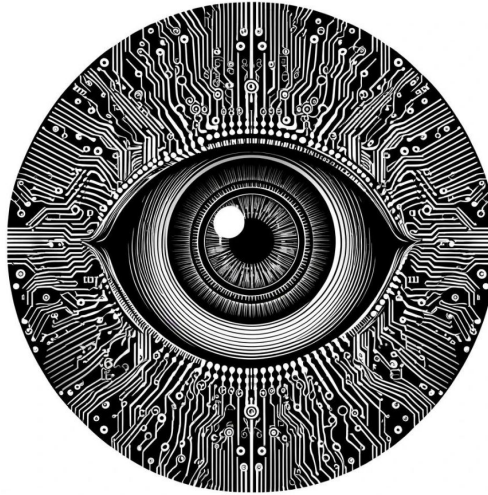
Agenda

- [Pooling in Neural Networks](#) (YouTube link)
- [Interpretability with Class Activation Mapping](#) (YouTube link)

Notebook

- [Interpretability with Class Activation Mapping](#) (Notebook on Google Colab)

Pooling in Neural Networks



Learning goals

- Explore image downsampling and reduction of network parameters with mean, max, and global pooling

Architecture of a convolutional network with pooling

```
nn.Sequential(  
    nn.Conv2d(3, 16, kernel_size=5, stride=1, padding=2),  
    nn.ReLU(),  
    nn.MaxPool2d(kernel_size=2, stride=2),  
  
    nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=2),  
    nn.ReLU(),  
    nn.MaxPool2d(kernel_size=2, stride=2),  
  
    nn.Flatten(),  
    nn.Linear(32 * 8 * 8, 120),  
    nn.ReLU(),  
    nn.Linear(120, 84),  
    nn.ReLU(),  
    nn.Linear(84, 10), # 10 classes, we are working with CIFAR 10  
)
```

airplane

automobile

bird

cat

deer

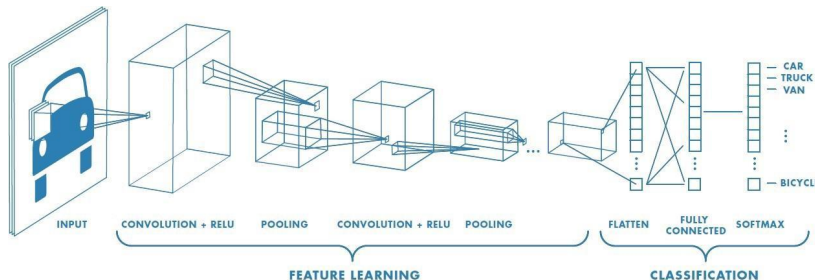
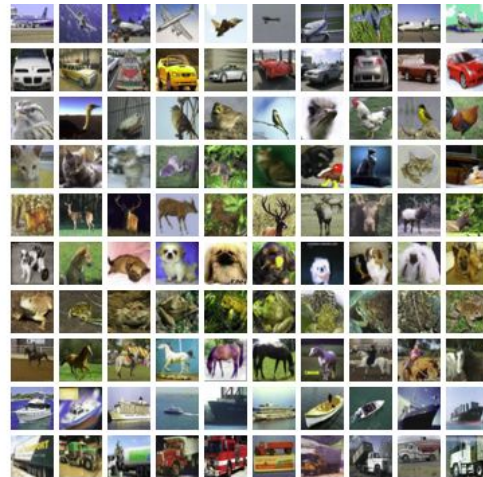
dog

frog

horse


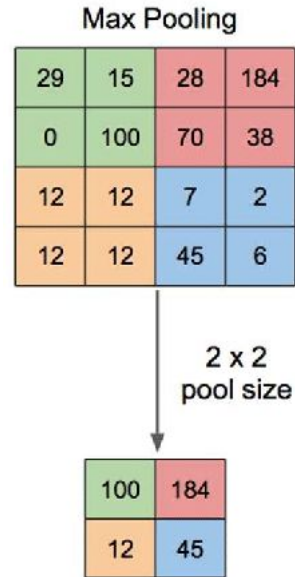
ship

truck



The CIFAR-10 dataset

Max Pooling



Mean Pooling

Average Pooling

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2
pool size

36	80
12	15

Input Image
Shape: 224x224



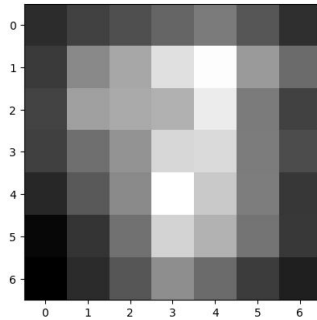
Mean Pooled Image (2x2 Kernel)
Shape: 112x112




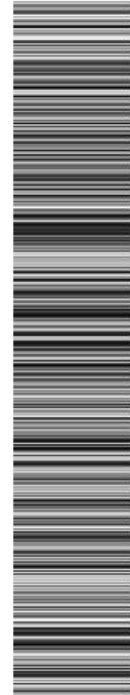
Global pooling aka adaptive pooling

```
import torch.nn as nn
```

```
class GlobalPooling(nn.Module):  
    def __init__(self):  
        super().__init__()  
        # Computes the average value of each activation map  
        self.avg_pool = nn.AdaptiveAvgPool2d((1, 1))  
        # Selects the max value of each activation map  
        self.max_pool = nn.AdaptiveMaxPool2d((1, 1))  
  
    def forward(self, x):  
        avg_pooled = self.avg_pool(x).flatten(1)  
        max_pooled = self.max_pool(x).flatten(1)  
        # Concatenates both poolings to produce a feature vector  
        return torch.cat([avg_pooled, max_pooled], dim=1)
```



Feature vector



Using global pooling to reduce parameters

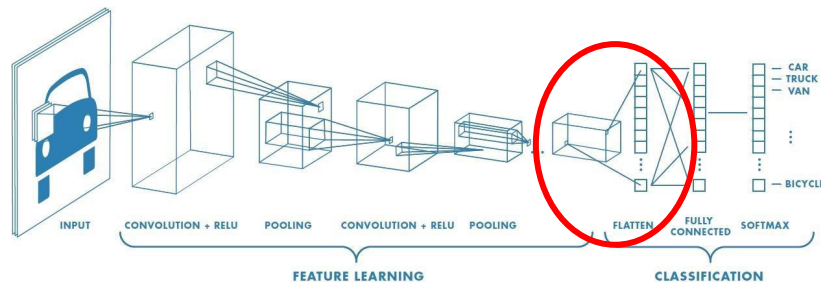
```
import torch
import torch.nn as nn

model = nn.Sequential(
    nn.Conv2d(3, 16, kernel_size=5, stride=1, padding=2),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),

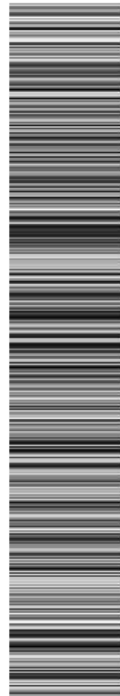
    nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=2),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),

    GlobalPooling(), # This replaces the Flatten layer

    nn.Linear(64, 120),
    nn.ReLU(),
    nn.Linear(120, 84),
    nn.ReLU(),
    nn.Linear(84, 10)
)
```



Feature vector



Summary

Pooling is a way to compress information

- Pooling allows us to do lossy compression while retaining important visual features
- Convolutions with stride > 1 achieve a similar effect at the cost of a higher number of parameters

Global pooling is an alternative to flattening activation maps

- We can create feature vectors (aka embeddings) using the global mean and/or max pooling operations

Further reading and references

A guide to convolution arithmetic for deep learning

- <https://arxiv.org/abs/1603.07285>

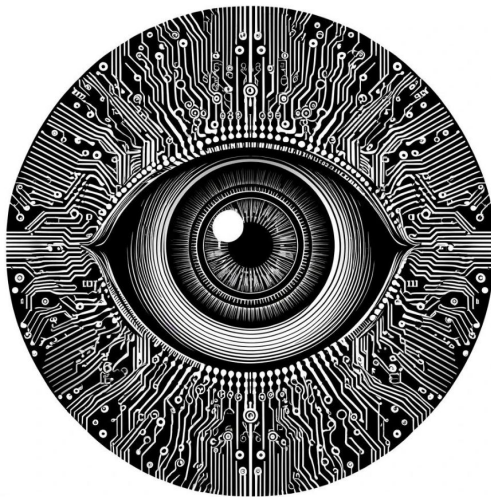
Network in network (1x1 convolutions)

- <https://arxiv.org/abs/1312.4400>

Hypercolumns for object segmentation and fine-grained localization

- https://openaccess.thecvf.com/content_cvpr_2015/papers/Hariharan_Hypercolumns_for_Object_2015_CVPR_paper.pdf

Interpretability with Class Activation Mapping



Learning goals

- Use class activation mapping to interpret the output of classifiers
- Describe tradeoffs between CAM and Grad-CAM

Interpreting a prediction



A pretrained Resnet34 with Imagenet1K_V1 weights says

“tabby, tabby cat” with

probability = 0.59

The resnet has 512 activations of shape $H = W = 7$ on its last layer, when we apply the 224x224 Resize on the input

Class Activation Mapping

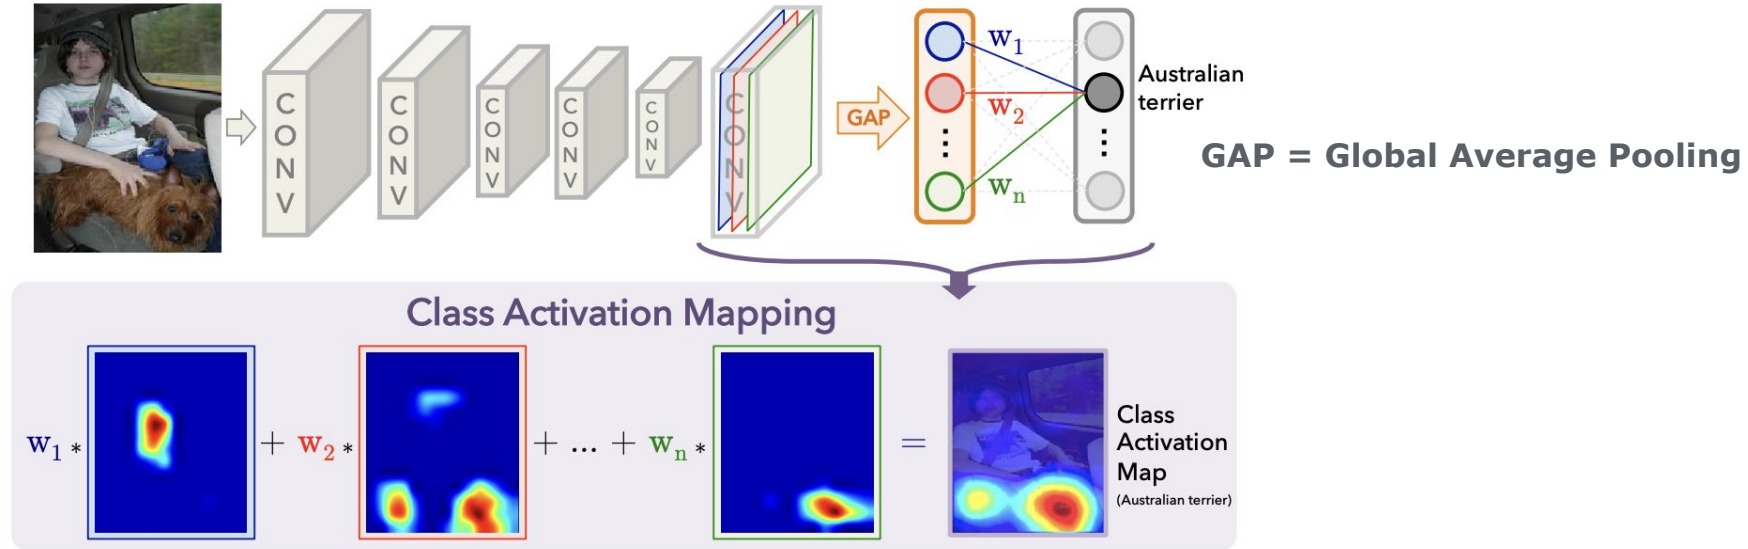


Figure 2. Class Activation Mapping: the predicted class score is mapped back to the previous convolutional layer to generate the class activation maps (CAMs). The CAM highlights the class-specific discriminative regions.

CAM with global average pooling

```
import torch
import torch.nn.functional as F
```

```
# Suppose fc_weights has shape [num_classes, channels] and activations has shape [1, batch_size, channels, H, W].
```

```
# We'll just show the relevant slices for the single 'class_idx' and the first image in the batch.
```

```
weight = fc_weights[class_idx]          # shape [channels]
```

```
act = activations[0][0]                  # shape [channels, H, W] - 7x7 in the case of resnet34
```

```
# -----
# 1) The "global average pooling" from a usual forward pass:
#    collapses (H, W) -> 1x1, giving us one value per channel.
# -----
```

```
pooled = F.adaptive_avg_pool2d(act.unsqueeze(0), 1) # shape [1, channels, 1, 1]
```

```
pooled = pooled.squeeze(0).squeeze(-1).squeeze(-1) # shape [channels]
```

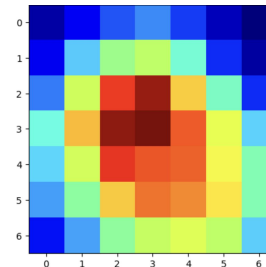
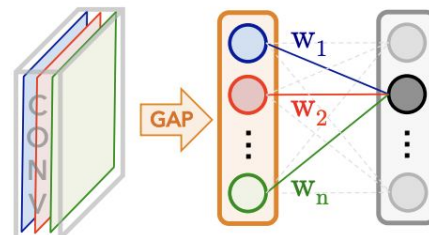
```
# 'pooled' is the channel-wise average. Multiplying by 'weight' then summing would give the final logit for 'class_idx'.
```

```
score = (pooled * weight).sum() # The single scalar logit for class_idx
```

```
# -----
# 2) Building the CAM:
#    multiply each channel map by its weight, then sum across channels.
# -----
```

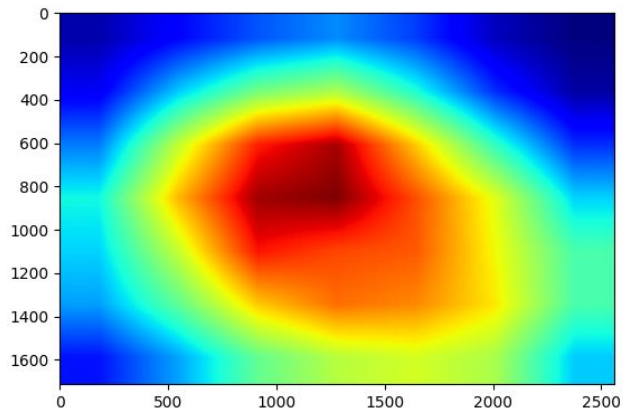
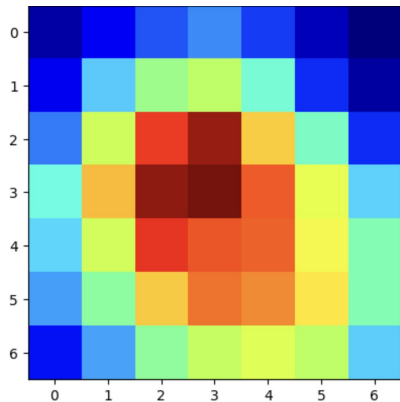
```
cam = (act * weight.view(-1, 1, 1)).sum(dim=0) # shape [H, W]
```

```
print(cam.shape) # [H, W] this is now shape 7x7
```



Resizing the map

```
cam_resized = np.array(Image.fromarray(cam).resize(img.size, resample=Image.BILINEAR))  
plt.imshow(cam_resized, cmap="jet");
```



Extracting the activations with a hook

```
# Hook for extracting the activations from the last convolutional layer
```

```
activations = []
```

```
def hook_fn(module, input, output):
```

```
    activations.append(output)
```

```
# Register the hook
```

```
layer_name = 'layer4' # Last convolutional block
```

```
hook = model._modules.get(layer_name).register_forward_hook(hook_fn)
```

```
# Forward pass
```

```
output = model(input_tensor)
```

```
# Remove the hook
```

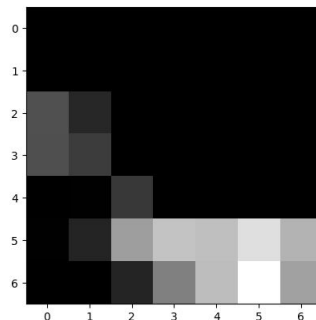
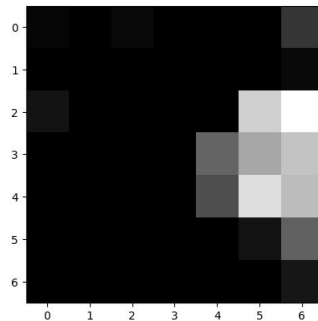
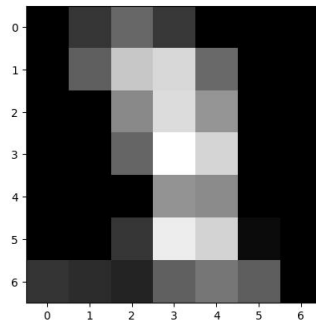
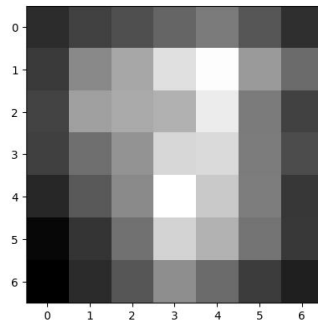
```
hook.remove()
```

```
# Get the weights of the fully connected layer
```

```
fc_weights = model.fc.weight.detach()
```

```
# Select the class index (e.g., 0 for 'tench')
```

```
class_idx = torch.argmax(output, dim=1).item()
```




Bilinear interpolation

$$w_x = \frac{x - x_1}{x_2 - x_1}$$


$$w_y = \frac{y - y_1}{y_2 - y_1}$$

$(x_1, y_1) = (0, 0)$ coordinates of Q_{11}

$(x_2, y_2) = (1, 1)$ coordinates of Q_{22}

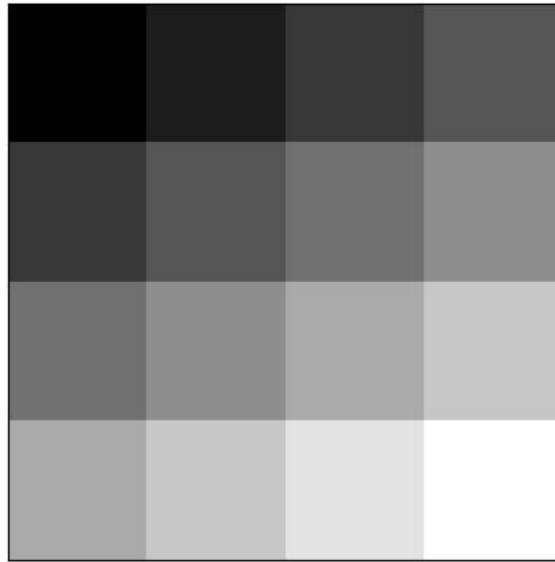
 $(x, y) = (0.6, 0.7)$ target point

input



$$Q = \begin{bmatrix} 10 & 20 \\ 30 & 40 \end{bmatrix}$$

output



$$f(x, y) = [1 - w_x \quad w_x]$$

$$\begin{bmatrix} Q_{11} & Q_{12} \\ Q_{21} & Q_{22} \end{bmatrix} \begin{bmatrix} 1 - w_y \\ w_y \end{bmatrix}$$

Bilinear interpolation in PyTorch

```
import torch
import torch.nn.functional as F
```

```
# Create a 1x1x2x2 tensor (batch_size x channels x height x width)
# Notice that the batch size and channel dimensions are created by wrapping
# the height and width tensor with two pairs of extra square brackets
```

```
input = torch.tensor([[[[10, 20],
                          [30, 40]]]],
                      dtype=torch.float32)
```

```
# Upscale to 4x4
```

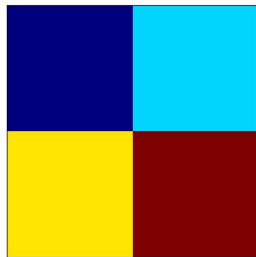
```
output = F.interpolate(input, size=(4, 4), mode='bilinear',
                       align_corners=True)
```

```
import matplotlib.pyplot as plt
```

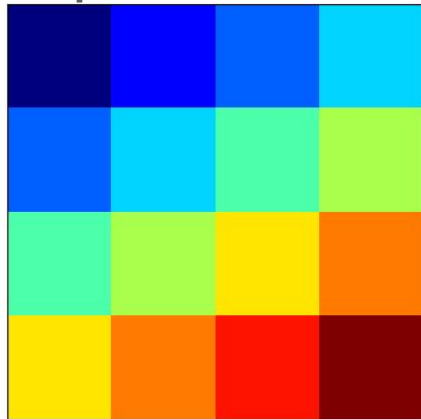
```
# The colormap is just for illustration of corner alignment
```

```
plt.imshow(input.squeeze(), cmap="jet")
plt.imshow(output.squeeze(), cmap="jet")
```

input (with colormap)
shape: 2x2



output (with colormap)
shape: 4x4



Overlaying the resized CAM on the image

```
# Compute CAM
# fc_weights[class_idx] is 1-dimensional, but the einsum equation 'ij' expects 2 dimensions.
# We need to unsqueeze to add a dimension to fc_weights[class_idx]
# The einsum operation here is a one liner that is equivalent to Global Average Pooling
cam = torch.einsum('ij,jkl->ikl', fc_weights[class_idx].unsqueeze( 0), activations[0][0])

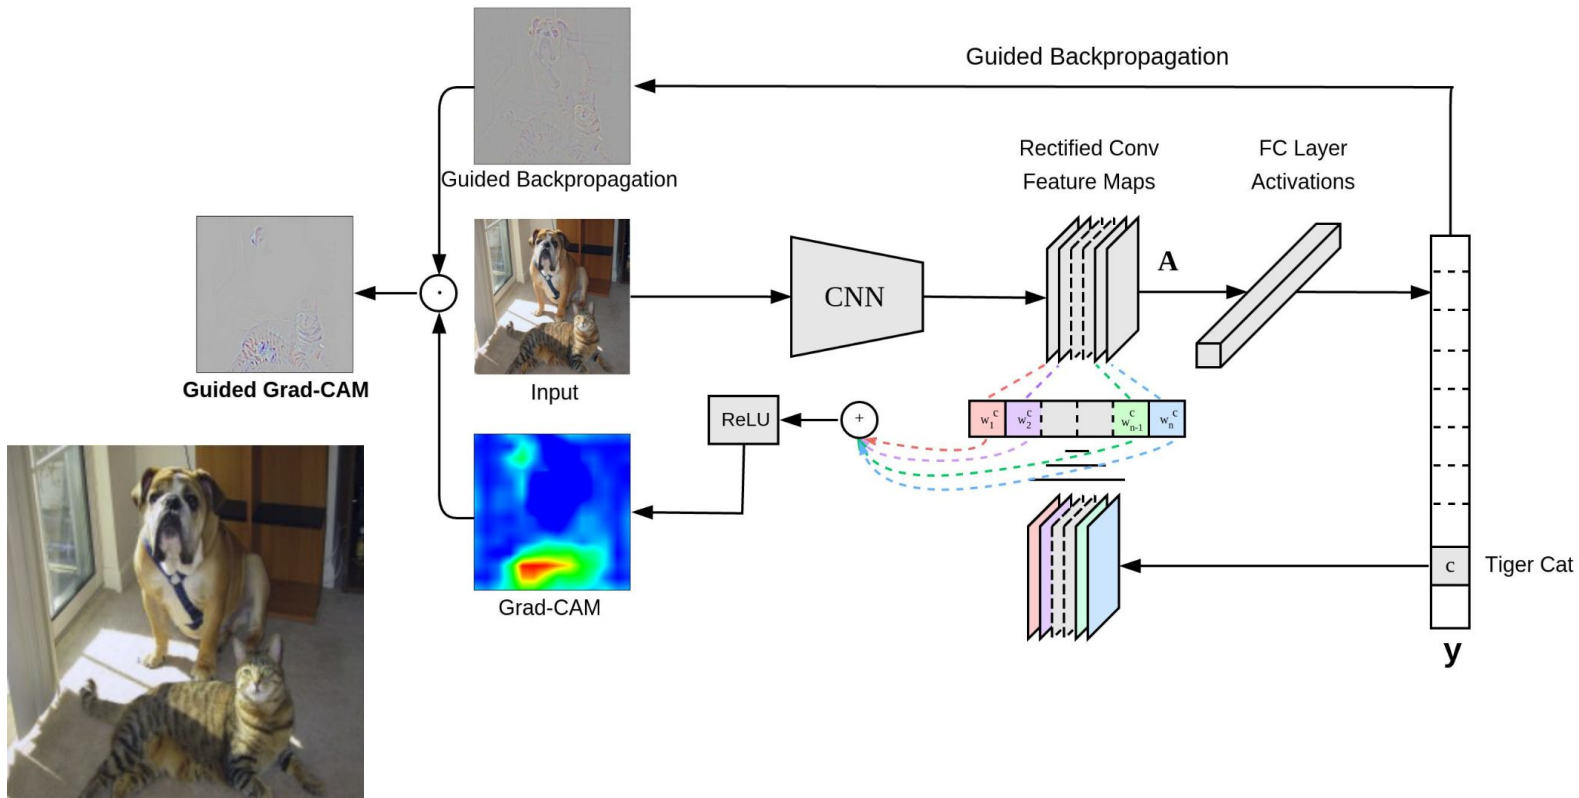
# Normalize CAM
cam = cam - cam.min()
cam = cam / cam.max()

# Resize CAM to match the input image
cam = cam.detach().numpy()
# Squeeze the cam array to remove the first dimension and convert to uint8
cam = cam.squeeze() # remove the first dimension
cam = (cam * 255).astype(np.uint8) # scale to 0-255 and convert to uint8
cam_resized = np.array(Image.fromarray(cam).resize(img.size, resample=Image.BICUBIC))

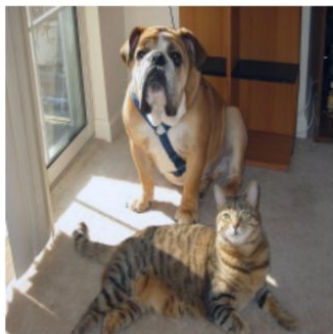
# Overlay CAM on the image
plt.imshow(img)
plt.imshow(cam_resized, cmap= 'jet', alpha=0.5)
plt.axis('off')
plt.show()
```



Guided Grad-CAM



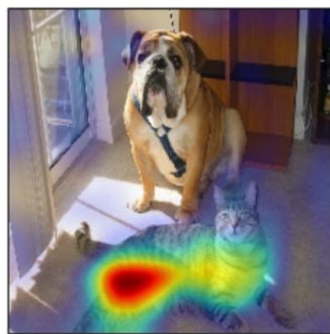
Combining Grad-CAM and Guided Backprop



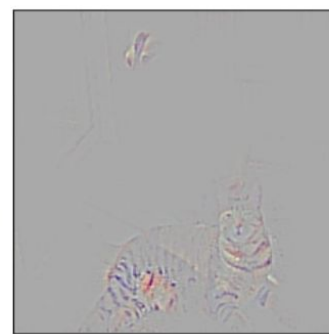
(a) Original Image



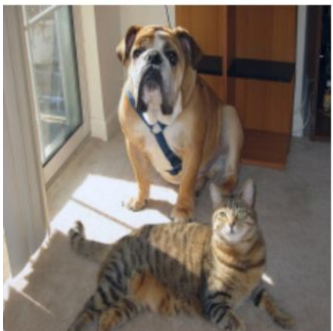
(b) Guided Backprop 'Cat'



(c) Grad-CAM 'Cat'



(d) Guided Grad-CAM 'Cat'



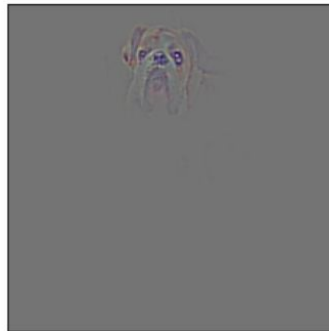
(g) Original Image



(h) Guided Backprop 'Dog'



(i) Grad-CAM 'Dog'



(j) Guided Grad-CAM 'Dog'

CAM vs GradCAM

- CAM needs a Global Average Pooling layer to be added to the model, GradCAM works with any architecture without changes
- GradCAM can visualize outputs of any layer, CAM is limited to the final layer
- CAM runs faster and requires less memory

Berlin 'lioness': Wild animal probably a boar, authorities say

21 July 2023

Share  Save 

Kathryn Armstrong
BBC News



Reuters

Image from [BBC News](#)

Michael Grubert, mayor of Kleinmachnow, said the spotted animal on the loose was most likely a boar

Lion or boar?




Image from [Berlin 'lioness' on loose 'is a wild boar' \(BBC News\)](#)

Summary

Class Activation Mapping (CAM) is a tool for explainability

- CAM helps us understand whether our decisions are well supported or based on spurious correlations
- CAM exposes hidden connections between inputs and decisions that affect model reliability and safety
- Grad-CAM allows us to extend the method to any network without changes to the architecture

CAM implementations steps

- Extract feature maps with hooks from final convolutional layer
- Project class weights onto activation maps
- Upsample and overlay heatmaps on input images

References

Learning Deep Features for Discriminative Localization

- https://openaccess.thecvf.com/content_cvpr_2016/html/Zhou_Learning_Deep_Features_CVPR_2016_paper.html

Class Activation Mapping explained

- https://github.com/fastai/fastbook/blob/master/18_CAM.ipynb

Basic guide to Numpy's einsum

- <https://ajcr.net/Basic-guide-to-einsum/>

Class activation mapping on fiftyone

- <https://voxel51.com/blog/exploring-gradcam-and-more-with-fiftyone/>

Resources

- [Github Repository](#)
- [YouTube playlist](#)
- [Discord channel](#)
#practical-computer-vision-workshops